

Hands-On Graphical Causal Modeling Using R

Johannes Textor

March 11, 2016

Graphs and paths

Model testing

Model equivalence

Causality theory

- ▶ A causality theory provides a language to encode causal relationships.
- ▶ A causality theory helps decide when, and how, causation can be inferred from domain knowledge and data.
- ▶ We use the theory based on **structural causal models**.

Some people with their own causality theories

Donald
Rubin



Judea
Pearl



Donald
Campbell



Phil
Dawid




Clive
Granger



*[...] all approaches to causation are variants or
abstractions of [...] structural theory [...].*
– Judea Pearl, “Causality”

Graphs and paths

- ▶ A **graph** consists of **nodes** (vertices) and **edges** (arrows or lines).



- ▶ We describe node relations using **kinship terminology**.
 - ▶ Z is a **parent** of E, D.
 - ▶ D is a **child** of Z, E.
 - ▶ D is a **descendant** of A.
 - ▶ A is an **ancestor** of D.

(terms for \leftrightarrow , $-$ (**spouse**, **neighbour**) are less used)

- ▶ A **path** is a sequence of connected nodes (moving against arrows is allowed), e.g. $Z \leftrightarrow A \rightarrow E \leftarrow Z \rightarrow D$.

The dagitty package

I am going to show how to work with structural causal models in R using the package 'dagitty'. The package is not yet on CRAN, so it needs to be installed as follows:

```
install.packages("devtools") #if you haven't already  
library(devtools)  
install_github("jtextor/dagitty/r")
```


If you wish to follow this talk by copy-pasting code (encouraged):
johannes-textor.name/leeds.pdf

Defining graphs in R

We define graphs using a textual syntax based on the program “graphviz”. A simple command exists to plot graphs.

```
g <- dagitty( "dag {  
a -> e  
e -> d  
z -> a  
z -> e  
z -> d  
}" )
```

```
plot(graphLayout(g))
```



Or, more briefly:

```
g <- dagitty( "dag{{a->e->d}<-z}" )
```

Defining graphs in R

The function `graphLayout` generates automatic layouts. This works OK, but of course you can also specify the layout yourself.

```
g <- dagitty( 'dag {  
  a [pos="0,0"]  
  e [pos="1,0"]  
  d [pos="2,0"]  
  z [pos="0,1"]  
  z->{a->e->d}}' )  
  
g <- dagitty( 'dag{  
  z->{a->e->d}}' )  
  
coordinates(g) <- list(  
  X=c(a=0,b=1,c=2,z=1),  
  Y=c(a=1,b=1,c=1,z=0))
```

Or build it in the `dagitty.net` GUI and download it.

```
g <- downloadGraph("dagitty.net/m331")
```


Ancestry relationships in R

The dagitty package contains functions for ancestral relationships.

```
parents( g, "Z" )
```

```
## [1] "W" "Y"
```

```
ancestors( g, "Z" )
```


```
## [1] "Z" "Y" "X" "W"
```

```
children( g, "W" )
```

```
## [1] "Y" "Z"
```

```
descendants( g, "W" )
```

```
## [1] "W" "Z" "T" "Y"
```




Paths in R

For detailed inspection of paths and their statuses, dagitty provides the function 'paths'. This function returns a list with two components, 'paths' and 'open'. I'll explain later what the 'open' is about.

```
paths( g, "Y", "Z" )$paths
```


```
## Y -> T <- Z  
## Y -> Z  
## Y <- W -> Z  
## Y <- X -> W -> Z
```



Structural causal models (informally)

- ▶ A structural causal model is a graph whose nodes represent variables, and whose edges represent causal relations.
- ▶ A **DAG** is a directed acyclic graph.
- ▶ We draw $X \rightarrow Y$ if X **might have** causal influence on Y . Otherwise we **omit** the arrow.
- ▶ We draw $X \leftrightarrow Y$ if there is a hidden variable influencing both X and Y .

Example



Z and A influence E

Something hidden influences A and Z

Z and E influence D

What do we mean by “influences”?

Structural causal models (formally)


A structural causal model consists of a graph $G = (V, E)$ and a set of functions $\{f_X\}$ such that, for each variable X :

$$X := f_X(\text{pa}_X, \epsilon_X)$$

- ▶ f_X is **any** deterministic function.
- ▶ pa_X is the set of all parents of X in G .
- ▶ ϵ_X is a random variable.

For $X \neq Y$, ϵ_X and ϵ_Y are **independent** unless $X \leftarrow \text{rightarrow} Y$.

Example



$$E := f_E(A, Z, \epsilon_E)$$

$$A := f_A(Z, \epsilon_A)$$


$$D := f_D(Z, E, \epsilon_D)$$

“Lady Nature looks up the values of Z and A when determining E ”

Structural equation models

- ▶ A **Structural Equation Model** (SEM, SEmodel) is a causal diagram where the functions are **restricted** to linear functions with additive Gaussian noise.

Example



$$E := \beta_{AD}A + \beta_{ZD}Z + \epsilon_E$$

$$A := \beta_{ZA}Z + \epsilon_A$$

$$D := \beta_{EA}E + \beta_{ZD}Z + \epsilon_D$$

Simulating data

It is often useful to simulate data from a DAG.

```
g <- "dag{a<->b<-c}"  
cor(simulateSEM(g,.5,.5))
```

```
##      a    b    c  
## a 1.0 0.5 0.0  
## b 0.5 1.0 0.5  
## c 0.0 0.5 1.0
```

```
g <- "dag{a<-b<->c}"  
cor(simulateSEM(g,.5,.5))
```

```
##      a    b    c  
## a 1.00 0.5 0.25  
## b 0.50 1.0 0.50  
## c 0.25 0.5 1.00
```

'simulateSEM' interprets the DAG as a structural equation model and samples path coefficients from the given interval.

Summary

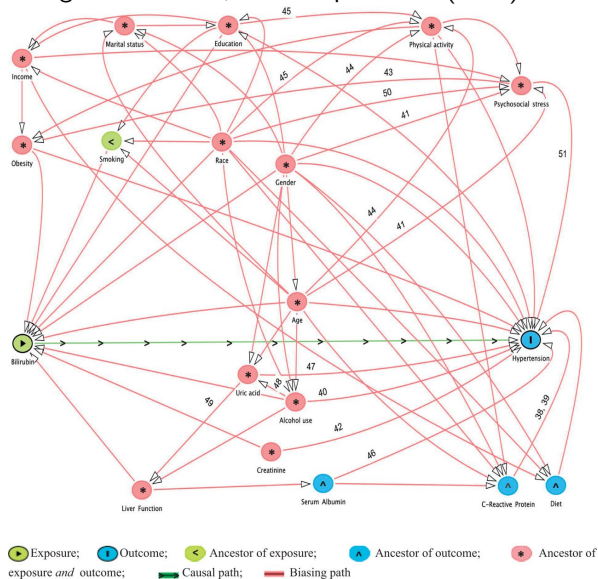
- ▶ **Structural causal models** consist of nodes (variables) and edges (functional relations).
- ▶ **DAGs** are models where all edges have arrowheads and that do not contain cycles.
- ▶ **Structural equation models** are DAGs where all edges represent **linear relations** with **additive Gaussian noise**.

Testing structural causal models

- ▶ Structural causal models are used for **inference** (e.g., computing covariate adjustment sets for regression).
- ▶ Such inferences depend on the validity of the model (frequent criticism: how do you know the model is correct?)
- ▶ Structural causal models can be **tested**.
- ▶ The key idea behind graphical model testing is **consistency**.



An example

Wang and Bautista, Int. J. Epidemiol. (2014)



Consistency

Given a probability distribution P and a graph G for some variables V , we say that P is **consistent** with G if there exists **any** set of functions and error terms by which G generates P .



Consistency is intimately related to how variables correlate. We will now formalize this intuition.

Conditional independence and vanishing covariance

A few definitions first:

- ▶ $\text{Cov}(X, Y \mid Z)$: partial covariance between X and Y given Z .
- ▶ $X \perp\!\!\!\perp Y$ means: X and Y are independent.
- ▶ $X \perp\!\!\!\perp Y \mid Z$ means: X and Y are independent given Z .
 Z and Y tell us no more about X than Z alone does.
- ▶ If $X \perp\!\!\!\perp Y$, then $\text{Cov}(X, Y) = 0$.
- ▶ If $X \perp\!\!\!\perp Y \mid Z$, then $\forall z : \text{Cov}(X, Y \mid Z = z) = 0$.

d -Separation

- ▶ A **collider** is a path of length 3 that looks like $X \rightarrow M \leftarrow Y$.
- ▶ All other paths of length 3 are called **non-colliders**:
 $X \rightarrow M \rightarrow Z$, $X \leftarrow M \leftarrow Y$, $X \leftarrow M \rightarrow Z$

d -separation

A set Z d -separates (**blocks**) a path, if


- ▶ The midpoint M of some non-collider is in Z ; or
- ▶ The midpoint M of some collider is not an ancestor of any variable in Z .

Theorem (Verma & Pearl, 1984)

*If all paths between X and Y are closed by the set Z , then $X \perp\!\!\!\perp Y \mid Z$ in **every** consistent probability distribution.*

Path inspection in R

```
g <- dagitty("dag{ X -> R -> S -> T <- U <- V -> Y  
T -> P }")
```



```
paths(g, "X", "Y")
```

```
## $paths
```

```
## [1] "X -> R -> S -> T <- U <- V -> Y"
```


```
##
```

```
## $open
```

```
## [1] FALSE
```

Path inspection in R

```
g <- dagitty("dag{ X -> R -> S -> T <- U <- V -> Y  
T -> P }")
```



```
paths(g, "X", "Y", "T")
```

```
## $paths
```

```
## [1] "X -> R -> S -> T <- U <- V -> Y"
```


```
##
```

```
## $open
```

```
## [1] TRUE
```

Path inspection in R

```
g <- dagitty("dag{ X -> R -> S -> T <- U <- V -> Y  
T -> P }")
```



```
paths(g, "X", "Y", "P")
```


```
## $paths  
## [1] "X -> R -> S -> T <- U <- V -> Y"  
##  
## $open  
## [1] TRUE
```

Verifying d-separation in R

```
g <- "dag { z -> m -> y } "  
impliedConditionalIndependencies(g)  
  
## y _||_ z | m  
  
confint( lm( y ~ z + m, data=simulateSEM( g ) ) )  
  
##           2.5 % 97.5 %  
## (Intercept) -0.045  0.104  
## z           -0.109  0.062  
## m           -0.558 -0.385
```


Discovering model misspecifications


d-Separation is a useful tool to check whether the DAGs that we postulate are in fact correct.



assumed model
(perfect measurement)

$$M_1 \perp\!\!\!\perp M_2 \mid X$$

$$X \perp\!\!\!\perp Y \mid \{M_1, M_2\}$$



true model
(measurement error)

$$M_1 \perp\!\!\!\perp M_2 \mid X$$

Model testing in R


```
# Simulate data from true model  
d <- simulateSEM("dag{X->{U1 M2}->Y U1->M1}", .6, .6)
```

```
# Test postulated model  
localTests( "dag{ X -> {M1 M2} -> Y }", d, "cis" )  
  
##               estimate std.error p.value  2.5% 97.5%  
## M1 _||_ M2 | X      -0.004      0.052 9.4e-01 -0.11 0.098  
## X _||_ Y | M1, M2    0.650      0.066 3.0e-21  0.52 0.779
```

Indeed, we find that the second implication is violated.

A real-world example


```
plot( getExample("Kampen") )
```



SSQ model of schizophrenic prodromal unfolding
van Kampen, European Psychiatry, 2014

A real-world example


```
tests <- localTests( getExample("Kampen"), d, "cis" )  
plotLocalTestResults( tests )
```



A real-world example

```
head(tests[order(tests$p.value),])
```

##		estimate	std.error	p.value	2.5%	97.5%
##	DET _ _ PER ALN	0.42	0.043	7.9e-22	0.340	0.51
##	ALN _ _ SAN AFF	0.22	0.024	2.2e-18	0.170	0.26
##	ALN _ _ CDR AFF	0.17	0.023	9.7e-14	0.129	0.22
##	CDR _ _ PER AFF	0.40	0.068	3.9e-09	0.270	0.54
##	PER _ _ SAN AFF	0.11	0.021	1.7e-07	0.069	0.15
##	CDR _ _ FTW AFF, AIS	0.16	0.031	6.0e-07	0.095	0.22



Summary

- ▶ Please test your models!
- ▶ 1 line of code using the daitty package ...

Equivalent models

- ▶ Sometimes two models have **exactly the same** testable implications.
- ▶ Such models are called **equivalent** – they cannot be distinguished by statistical means alone.

```
impliedConditionalIndependencies("dag{ x -> m -> y }")
```


```
## x _||_ y | m
```

```
impliedConditionalIndependencies("dag{ x <- m <- y }")
```


```
## x _||_ y | m
```

Markov equivalence

- Models with the same testable implications are called **Markov equivalent**.





$$D \perp\!\!\!\perp A \ ; \ A \perp\!\!\!\perp Z$$



$$D \perp\!\!\!\perp A \ ; \ A \perp\!\!\!\perp Z$$


- Two models are Markov equivalent iff they have the same “skeleton” (edges w/o arrowheads) and the same “immoralities” (children of “unmarried” = unlinked parents).




Equivalent models in R

```
g <- dagitty( "dag{  
  B->{A Z}->X->W->Y  
  C->{Z D}->Y}" )
```

```
plot(g)
```




```
for( g2 in equivalentDAGs(g) ){  
  plot(g2)  
}
```



Equivalent models


```
plot(equivalenceClass(g))
```

The equivalence class is a graph that contains both arrows and lines (undirected edges). The lines can be oriented in either direction, as long as **no immorality** and **no cycle** is created.




Equivalence classes

How many equivalent models do you think this DAG has?



Equivalence classes

How many equivalent models do you think this DAG has?

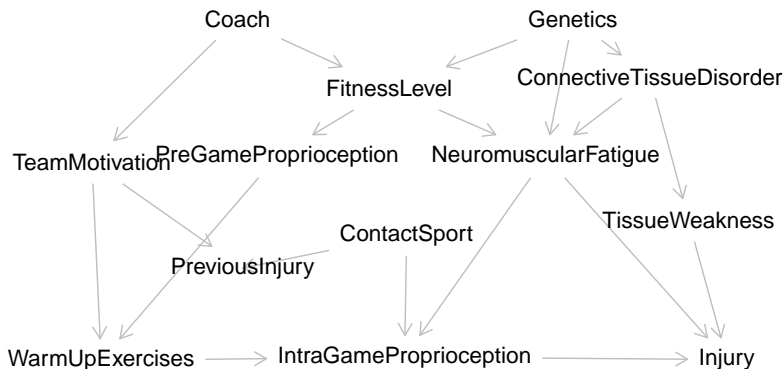


```
length(equivalentDAGs( g ))
```

```
## [1] 1
```

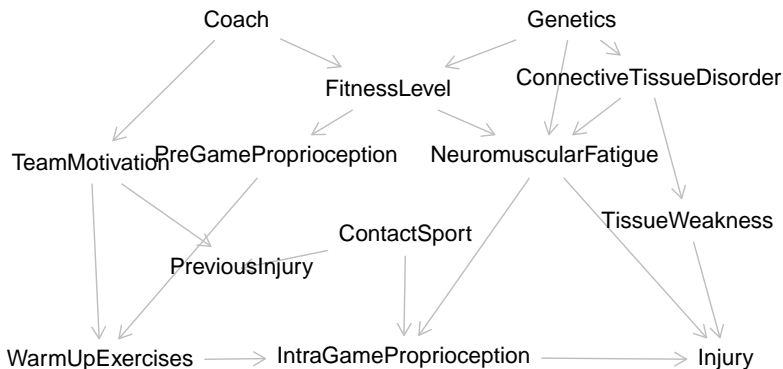
Equivalence classes

How many equivalent models do you think this DAG has?



Equivalence classes

How many equivalent models do you think this DAG has?

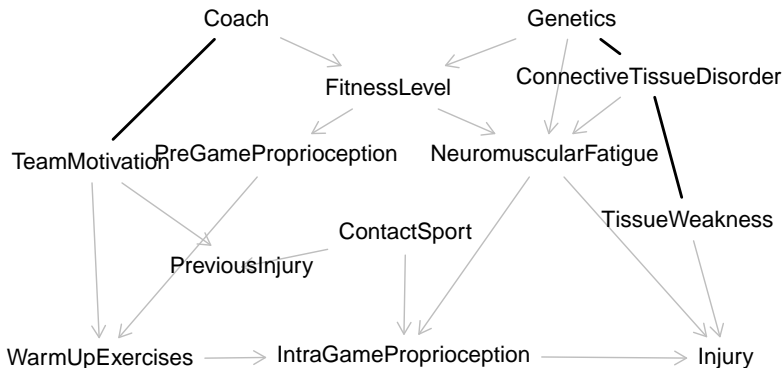


```
length(equivalentDAGs( g ))
```

```
## [1] 6
```

Equivalence classes

```
plot(equivalenceClass(g))
```




Only 3 edges can be reversed!

Summary

- ▶ Markov equivalent models are statistically indistinguishable (unless parametric assumptions are made).
- ▶ **Immoralities** are key to break Markov equivalence.
- ▶ There can be surprisingly few equivalent model if there are many immoralities.

Conclusion

- ▶ dagitty is known as a web-based user interface for drawing and analyzing graphs.
- ▶ An R package is now available that does everything the GUI can, and much more.
- ▶ A dagitty-based companion for Judea Pearl's new book is available at dagitty.net/primer.



Acknowledgements

- ▶ George Ellison and Mark Gilthorpe, Leeds :)
- ▶ Bela van der Zander and Maciej Liśkiewicz, Lübeck, Germany
- ▶ Felix Thoemmes, Cornell
- ▶ Judea Pearl and Andrew Forney, UCLA